

model-selection-in-ml-with-python

April 29, 2024

```
[ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
```

Importing the dependencies

```
[ ]: # importing the models
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
```

We will be working on the Heart Disease dataset

```
[ ]: # loading the csv data to a Pandas DataFrame
heart_data = pd.read_csv('/content/heart.csv')
```

```
[ ]: # print first 5 rows of the dataset
heart_data.head()
```

```
[ ]: 
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[ ]: # number of rows and columns in the dataset
heart_data.shape
```

```
[ ]: (303, 14)
```

```
[ ]: # checking for missing values
heart_data.isnull().sum()
```

```
[ ]: age      0
     sex      0
     cp       0
     trestbps  0
     chol     0
     fbs      0
     restecg  0
     thalach  0
     exang    0
     oldpeak  0
     slope    0
     ca       0
     thal     0
     target   0
     dtype: int64
```

```
[ ]: # checking the distribution of Target Variable
heart_data['target'].value_counts()
```

```
[ ]: 1    165
     0    138
     Name: target, dtype: int64
```

1 -> Defective Heart

0 -> Healthy Heart

Splitting the Features and Target

```
[ ]: X = heart_data.drop(columns='target', axis=1)
     Y = heart_data['target']
```

```
[ ]: print(X)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	
3	56	1	1	120	236	0	1	178	0	0.8	
4	57	0	0	120	354	0	1	163	1	0.6	
..	
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	

301	57	1	0	130	131	0	1	115	1	1.2
302	57	0	1	130	236	0	0	174	0	0.0

	slope	ca	thal
0	0	0	1
1	0	0	2
2	2	0	2
3	2	0	2
4	2	0	2
..
298	1	0	3
299	1	0	3
300	1	2	3
301	1	1	3
302	1	1	2

[303 rows x 13 columns]

```
[ ]: print(Y)
```

```
0      1
1      1
2      1
3      1
4      1
..
298    0
299    0
300    0
301    0
302    0
```

Name: target, Length: 303, dtype: int64

```
[ ]: X = np.asarray(X)
      Y = np.asarray(Y)
```

Model Selection

1. Comparing the models with default hyperparameter values using Cross Validation

```
[ ]: # list of models
models = [LogisticRegression(max_iter=1000), SVC(kernel='linear'),
          ↪KNeighborsClassifier(), RandomForestClassifier(random_state=0)]
```

```
[ ]: def compare_models_cross_validation():
      for model in models:
```

```

cv_score = cross_val_score(model, X, Y, cv=5)
mean_accuracy = sum(cv_score)/len(cv_score)
mean_accuracy = mean_accuracy*100
mean_accuracy = round(mean_accuracy, 2)

print('Cross Validation accuracies for the',model,'=', cv_score)
print('Accuracy score of the ',model,'=',mean_accuracy,'%')
print('-----')

```

```
[ ]: compare_models_cross_validation()
```

```

Cross Validation accuracies for the LogisticRegression(max_iter=1000) =
[0.80327869 0.8852459 0.85245902 0.86666667 0.75      ]
Accuracy score of the  LogisticRegression(max_iter=1000) = 83.15 %
-----
Cross Validation accuracies for the SVC(kernel='linear') = [0.81967213 0.8852459
0.80327869 0.86666667 0.76666667]
Accuracy score of the  SVC(kernel='linear') = 82.83 %
-----
Cross Validation accuracies for the KNeighborsClassifier() = [0.60655738
0.6557377 0.57377049 0.73333333 0.65      ]
Accuracy score of the  KNeighborsClassifier() = 64.39 %
-----
Cross Validation accuracies for the RandomForestClassifier(random_state=0) =
[0.85245902 0.90163934 0.81967213 0.81666667 0.8      ]
Accuracy score of the  RandomForestClassifier(random_state=0) = 83.81 %
-----

```

Inference: For the Heart Disease dataset, **Random Forest Classifier** has the Highest accuracy value with default hyperparameter values

2. Comparing the models with different Hyperparameter values using GridSearchCV

```
[ ]: # list of models
models_list = [LogisticRegression(max_iter=10000), SVC(),
↳KNeighborsClassifier(), RandomForestClassifier(random_state=0)]
```

```
[ ]: # creating a dictionary that contains hyperparameter values for the above
↳mentioned models
```

```

model_hyperparameters = {

    'log_reg_hyperparameters': {

        'C' : [1,5,10,20]

    },

```

```

    'svc_hyperparameters': {

        'kernel' : ['linear', 'poly', 'rbf', 'sigmoid'],
        'C' : [1,5,10,20]
    },

    'KNN_hyperparameters' : {

        'n_neighbors' : [3,5,10]
    },

    'random_forest_hyperparameters' : {

        'n_estimators' : [10, 20, 50, 100]
    }
}

```

```
[ ]: type(model_hyperparameters)
```

```
[ ]: dict
```

```
[ ]: print(model_hyperparameters.keys())
```

```
dict_keys(['log_reg_hyperparameters', 'svc_hyperparameters',
'KNN_hyperparameters', 'random_forest_hyperparameters'])
```

```
[ ]: model_hyperparameters['log_reg_hyperparameters']
```

```
[ ]: {'C': [1, 5, 10, 20]}
```

```
[ ]: model_keys = list(model_hyperparameters.keys())
print(model_keys)
```

```
['log_reg_hyperparameters', 'svc_hyperparameters', 'KNN_hyperparameters',
'random_forest_hyperparameters']
```

```
[ ]: model_keys[0]
```

```
[ ]: 'log_reg_hyperparameters'
```

```
[ ]: model_hyperparameters[model_keys[0]]
```

```
[ ]: {'C': [1, 5, 10, 20]}
```

Applying GridSearchCV

```
[ ]: def ModelSelection(list_of_models, hyperparameters_dictionary):

    result = []

    i = 0

    for model in list_of_models:

        key = model_keys[i]

        params = hyperparameters_dictionary[key]

        i += 1

        print(model)
        print(params)
        print('-----')

        classifier = GridSearchCV(model, params, cv=5)

        # fitting the data to classifier
        classifier.fit(X,Y)

        result.append({
            'model used' : model,
            'highest score' : classifier.best_score_,
            'best hyperparameters' : classifier.best_params_
        })

        result_dataframe = pd.DataFrame(result, columns = ['model used', 'highest_
        ↪score', 'best hyperparameters'])

    return result_dataframe
```

```
[ ]: ModelSelection(models_list, model_hyperparameters)
```

```
LogisticRegression(max_iter=10000)
{'C': [1, 5, 10, 20]}
-----
SVC()
{'kernel': ['linear', 'poly', 'rbf', 'sigmoid'], 'C': [1, 5, 10, 20]}
-----
KNeighborsClassifier()
{'n_neighbors': [3, 5, 10]}
-----
RandomForestClassifier(random_state=0)
```

```
{'n_estimators': [10, 20, 50, 100]}
```

```
[ ]:          model used  highest score  \
0      LogisticRegression(max_iter=10000)    0.831585
1                               SVC()         0.828306
2          KNeighborsClassifier()           0.643880
3  RandomForestClassifier(random_state=0)     0.838087
```

```
          best hyperparameters
0                {'C': 5}
1  {'C': 1, 'kernel': 'linear'}
2                {'n_neighbors': 5}
3                {'n_estimators': 100}
```

Random Forest Classifier with `n_estimators = 100` has the highest accuracy

```
[ ]:
```